

Line Coverage Improvement of Objdump using AFL Fuzzing and Neural Networks

By Yuval Schaal and Ido Michael

Motivation

- As computer engineering students, we are interested in machine learning.
- We wanted to use machine learning to dive into software reliability.
- Test if there was a way to make software more reliable using ML.



Project layout

1. Learning the basics and documentation of AFL fuzzer, instrument into a simple hello world.
2. Generate another hello world with Gcov & Lcov instrumentations and fuzz hello world app to generate seeds.
3. Work on a real world program such as binutils and generate training/testing data for the network (fuzzing seeds), Fuzzing compiled objects with a simple small input to mutate it and get better coverage.
4. Train the Neural Network with this data.
5. Testing the code coverage with Lcov, to see the unique crashes and line coverage.

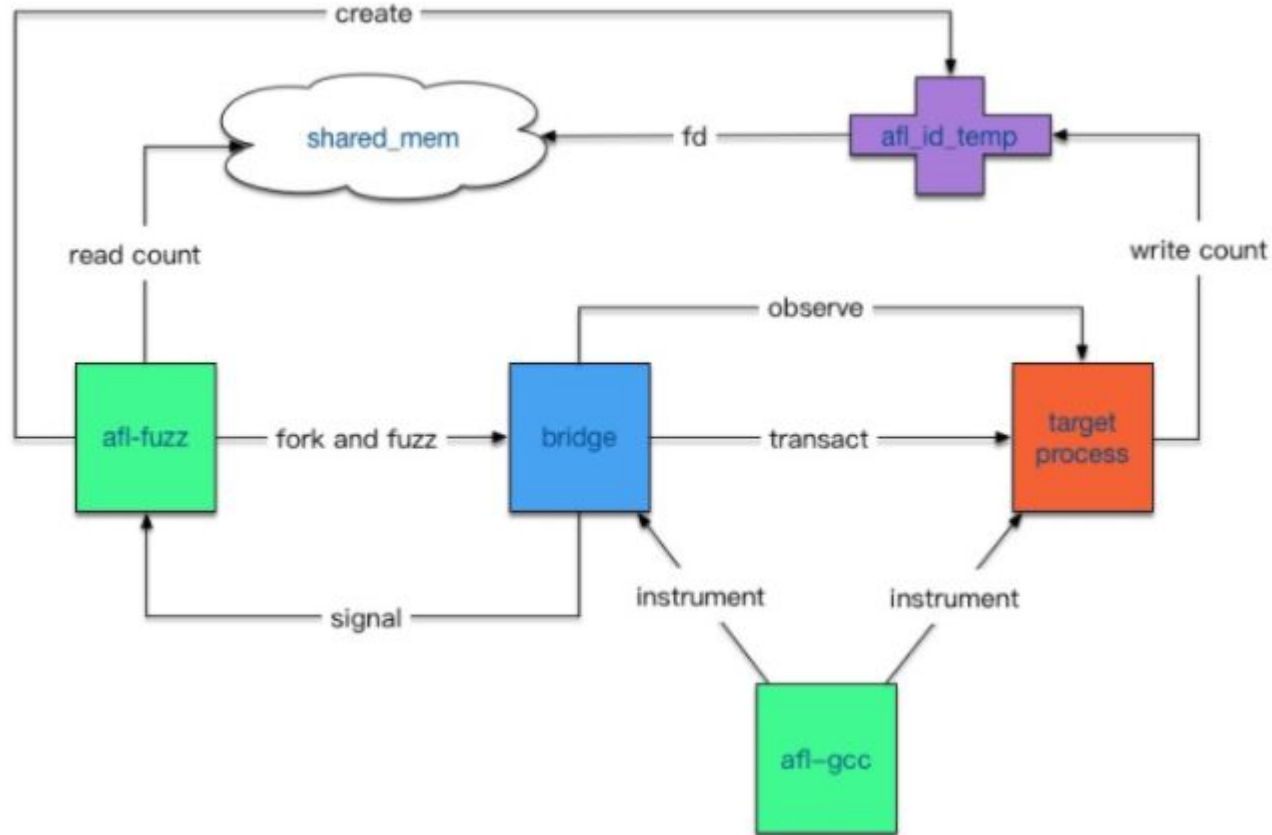
Previous Work

- Our mentor for this project is Dongdong She, a current Columbia CS PHD student of Prof. Junfeng.
- His work, very recently done, was to run the coverage project on readelf, which displays information about one or more ELF format object files, and see how to get better line coverage for software reliability.

Introduction to AFL Fuzzing

1. Compiling a binary with AFL instrumentations
2. Fuzzing those compiled objects with a simple small input to mutate it, all to get a better code coverage.
3. Testing the code coverage with Lcov, to see the unique crashes and line coverage.
4. AFL also contains a tool called afl-showmap, we will review it during the bitmap generation.

AFL Fuzzing diagram



Fuzzing process

- The AFL has a file called afl-gcc: this is the adapter for gcc compiler that creates the instrumentation into the obj file.
- After instrumenting, we need to fuzz the compiled file with some input and flag (depends on the program). This fuzzing is mutating the input file to get to new functions and get more coverage.
- We ran 8 AFL threads which synchronise once in a while, we got a total of about 6K seeds as training data and 8K for testing, total of 14K seeds.

How it actually looks like?

american fuzzy lop 1.86b (test)

process timing		overall results
run time : 0 days, 0 hrs, 0 min, 2 sec		cycles done : 0
last new path : none seen yet		total paths : 1
last uniq crash : 0 days, 0 hrs, 0 min, 2 sec		uniq crashes : 1
last uniq hang : none seen yet		uniq hangs : 0
cycle progress	map coverage	
now processing : 0 (0.00%)	map density : 2 (0.00%)	
paths timed out : 0 (0.00%)	count coverage : 1.00 bits/tuple	
stage progress	findings in depth	
now trying : havoc	favorable paths : 1 (100.00%)	
stage execs : 1464/5000 (29.28%)	new edges on : 1 (100.00%)	
total execs : 1697	total crashes : 39 (1 unique)	
exec speed : 626.5/sec	total hangs : 0 (0 unique)	
fuzzing strategy yields	path geometry	
bit flips : 0/16, 1/15, 0/13	levels : 1	
byte flips : 0/2, 0/1, 0/0	pending : 1	
arithmetics : 0/112, 0/25, 0/0	pend fav : 1	
known ints : 0/10, 0/28, 0/0	own finds : 0	

Background Objdump

- After Dongdong had successful results with Readelf, we decided to test a smaller, similar file: Objdump, also a part of the GNU Binary Utilities.
- Objdump is a program for displaying various information about object files on Unix-like systems. For instance, it can be used as a disassembler to view an executable in assembly form.
- During the process we have also used this tool to check if the compiled object files contained gcov/afl instrumentations.
- There are multiple flags you can use with Objdump, -D causes the highest function coverage which is why we chose to use it.

How Objdump Looks Like?

```
[root@linux-server root]#  
[root@linux-server root]# objdump -D a.out | grep -A20 main.:  
08048460 <main>:  
  8048460:    55                push    %ebp  
  8048461:    89 e5             mov     %esp,%ebp  
  8048463:    83 ec 08          sub     $0x8,%esp  
  8048466:    90                nop  
  8048467:    c7 45 fc 00 00 00 00 movl    $0x0,0xfffffffffc(%ebp)  
  804846e:    89 f6             mov     %esi,%esi  
  8048470:    83 7d fc 09       cmpl    $0x9,0xfffffffffc(%ebp)  
  8048474:    7e 02             jle     8048478 <main+0x18>  
  8048476:    eb 18             jmp     8048490 <main+0x30>  
  8048478:    83 ec 0c          sub     $0xc,%esp  
  804847b:    68 08 85 04 08    push    $0x8048508  
  8048480:    e8 93 fe ff ff    call    8048318 <_init+0x38>  
  8048485:    83 c4 10          add     $0x10,%esp  
  8048488:    8d 45 fc          lea     0xfffffffffc(%ebp),%eax  
  804848b:    ff 00             incl    (%eax)  
  804848d:    eb e1             jmp     8048470 <main+0x10>  
  804848f:    90                nop  
  8048490:    b8 00 00 00 00    mov     $0x0,%eax  
  8048495:    c9                leave  
  8048496:    c3                ret  
[root@linux-server root]# _
```

Background - Gcov

- A source code coverage analysis and statement-by-statement profiling tool
- Generates exact counts of the number of times each statement in a program is executed and annotates source code to add instrumentation

```
#include <stdio.h>

int
main (void)
{
    int i;

    for (i = 1; i < 10; i++)
    {
        if (i % 3 == 0)
            printf ("%d is divisible by 3\n", i);
        if (i % 11 == 0)
            printf ("%d is divisible by 11\n", i);
    }

    return 0;
}
```

C program before running gcov

```
#include <stdio.h>

int
main (void)
{
1   int i;

10  for (i = 1; i < 10; i++)
    {
9      if (i % 3 == 0)
3          printf ("%d is divisible by 3\n", i);
9      if (i % 11 == 0)
#####          printf ("%d is divisible by 11\n", i);
9    }

1    return 0;
1 }
```


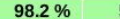

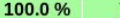

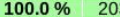

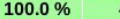

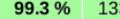

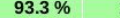
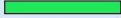
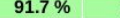
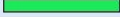
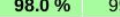
Gcov file

Background - Lcov

1. LCOV is a graphical front-end for GCC's coverage testing tool gcov.
2. After generating all of the necessary files we needed to run Lcov to create a HTML summary of the code coverage and the actual crashes, unique functions etc.


LCOV - code coverage report

Current view: top level		Hit		Total	
Test: libbash test coverage		Lines:	20640	34749	
Date: 2011-05-26		Functions:	1184	1287	
		Branches:	15689	37086	

Directory ▾	Line Coverage		Functions ▾	
src/core		95.7 % 314 / 328		98.2 % 55 / 56
test		97.0 % 98 / 101		100.0 % 72 / 72
src/builtins/tests		98.6 % 144 / 146		100.0 % 203 / 203
src/builtins		98.6 % 214 / 217		100.0 % 45 / 45
src/core/tests		98.9 % 351 / 355		99.3 % 133 / 134
./src/builtins		100.0 % 9 / 9		93.3 % 14 / 15
src		100.0 % 35 / 35		91.7 % 11 / 12
./src/core		100.0 % 190 / 190		98.0 % 99 / 101

Background - Keras and Tensorflow

- Tensorflow is an open source software library, for this project we used it for the machine learning applications utilizing neural networks.
- Keras runs on top of Tensorflow and is designed to enable fast experimentation with deep neural networks.



Environment

- Google Cloud Platform
- Compute Engine VM Instance
- GPU NVIDIA Tesla K80
- 8 Core CPU
- 500 GB Standard persistent disk
- Ubuntu 14.04



Google Cloud Platform

Generating the bitmap data


- As we mentioned the AFL has a tool called showmap.
- This tool runs the targeted binary and displays the contents of the trace bitmap in a human-readable form It's very useful in scripts to eliminate redundant inputs and perform other checks.
- We have fed to this tool the compiled binary with the mutated seeds that we generated, to create bitmaps (our NN data).

Cleaning the data

- We removed any seeds which was larger than 10K.
- Then we needed to put all the unique bitmaps of both the training and testing (fuzzer01, fuzzer02) into a single file:
 - `awk '{if (!a[$0]++) print}' ./fuzzer01/bitmaps/* > tempfile`
 - `awk '{if (!a[$0]++) print}' ./fuzzer02/bitmaps/* >> tempfile`
- After that we used a python script: `convert_bitmap.py` to create both the `convert_bitmaps` folder, as well as the `convert_testing_bitmaps`.

Training the model

- To train the model we used four layers in the neural network with 70 epochs, and 144 steps per each epoch.
- The results of this model had 97.25% accuracy.



Running gcov on the NN output

- So far, we got up to this step and we have the outputs.
- After running the NN and Mutate_gcov we get millions of mutated seeds.
- This step includes running those seeds on the gcov files one by one to create coverage report.
- With this report we can see if we got better results using the NN.
- Lcov will generate a html file that will be a report of the line coverage for objdump

Lcov results



LCOV - code coverage report

Current view: [top level](#) - /home/reliableProj/ReliableProj/afl_gcov/binutils-2.30/binutils

Test: coverage1.info

Date: 2018-04-22 05:31:48

	Hit	Total	Coverage
Lines:	83	249	33.3 %
Functions:	7	15	46.7 %

Filename	Line Coverage	Functions
size.c	<div><div></div></div> 33.3 % 83 / 249	46.7 % 7 / 15

Generated by: [LCOV version 1.13-14-ga5dd952](#)

Conclusions

1. Fuzzing is a really powerful tool on it's own.
2. Fuzzing and Neural Networks can shorten and even eliminate the current debugging process of software development.
3. All of the pipeline of fuzzing the seeds to generate data, creating bitmaps out of those seeds, training the NN, feeding millions of seeds to Gcov, has a really high complexity, it took us about 3 days to run end-to-end with a powerful GPU.

Future work

1. There are many improvement and changes that can be made, running on different programs, changing the data amounts or parameters, trying different network architectures.
2. A continuous, automated process can be achieved with LSTMs (Long-short term memory), training a model and once in a while feeding it with new seeds that the fuzzer has generated.
3. Because of the heavy complexity of the pipeline we didn't really meet a limit of the fuzzer/the network. An interesting experiment would be to test the limits of the system and see when the effectiveness drops.

References

Rawat, Sanjay, and Vivek Jain. VUzzer: Application-Aware Evolutionary Fuzzing.
(https://www.cs.vu.nl/~herbertb/download/papers/vuzzer_ndss17.pdf)

Prashant Anantharaman, Michael C. Millian, Sergey Bratus, Meredith L. Patterson, "Input Handling Done Right: Building Hardened Parsers Using Language-Theoretic Security", Cybersecurity Development (SecDev) 2017 IEEE, pp. 4-5, 2017
(<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=7839812>)

Rick, et al. "Complementing Model Learning with Mutation-Based Fuzzing." [1611.02429] Complementing Model Learning with Mutation-Based Fuzzing, 8 Nov. 2016, arxiv.org/abs/1611.02429.
(<https://arxiv.org/abs/1611.02429>)

Zhengzi Xu, Bihuan Chen, Mahinthan Chandramohan, Yang Liu, Fu Song, "SPAIN: security patch analysis for binaries towards understanding the pain and pills", Software Engineering Proceedings of the 39th International Conference (<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=7958599>)

Thanks!

Any questions?

You can find us at

im2492@columbia.edu & ys3055@columbia.edu